



Machine Learning Fundamentals

Assignment: 1.4 AI Lab

Professor:
Dr. Asabe Dawudu

Submitted On:
October 29, 2025

Submitted By:
Sandeep Sudesh Kumar

Student ID: 2927496

AI Lab

Section 1 - Practice with a Large Language Model (LLM)

Tool Used

ChatGPT (GPT-5)

Context

For this exercise, I explored how generative AI responds to structured prompting. My goal was to understand how an LLM interprets context, constraints, and tone when the same prompt is tested in different chats.

Prompt Framework

Prompt:

“You are an AI assistant helping a university student design a personal time-management plan for online courses. Create a one-week schedule balancing study, work, and rest. The plan should be concise and motivational.”

Framework applied:

- **Context:** Student managing studies and part-time work.
- **Task outline:** Generate a realistic 7-day schedule.
- **Constraints:** Keep it concise (under 200 words) and include motivation or reflection.

Test 1 - Default Chat

Prompt Used:

“You are an AI assistant helping a university student design a personal time-management plan for online courses. Create a one-week schedule balancing study, work, and rest. The plan should be concise and motivational.”

Output Summary:

The model generated a clear **Weekly Time-Management Plan** with detailed daily activities from Monday to Sunday, including study sessions, personal development, and rest periods.

Observation:

The plan was practical and well-structured but somewhat generic, lacking personalization. The motivational tone was mild but effective for general audiences.

Assessment: 7 / 10



Screenshot:

You are an AI assistant helping a university student design a personal time-management plan for online courses. Create a one-week schedule balancing study, work, and rest. The plan should be concise and motivational.

Here's a concise and motivational one-week time-management plan designed for an online university student balancing study, work, and rest:

Weekly Time-Management Plan

Goal:

Stay consistent, manage time wisely, and maintain a healthy balance between learning, working, and personal well-being.

Monday–Friday

Time	Activity	Purpose
7:00–8:00 AM	Morning routine (stretch, breakfast, quick journaling)	Start the day calm and focused
8:00–12:00 PM	Work or Classes	Focus on key academic/work priorities
12:00–1:00 PM	Lunch + short walk	Refresh your mind
1:00–3:00 PM	Study session (assignments, readings, or videos)	Deep focus on coursework
3:00–4:00 PM	Short break or power nap	Recharge energy
4:00–7:00 PM	Work shift / practical tasks	Productive grind time
7:00–8:00 PM	Dinner + light exercise	Balance body and mind
8:00–10:00 PM	Study review or reflection	Strengthen memory, plan next day
10:30 PM	Wind down (music, meditation, or light reading)	Peaceful end to the day

Figure 1. Test 1 – Default Chat Output (Weekly Time-Management Plan, Part 1)

Saturday

Time	Activity	Purpose
9:00–10:00 AM	Morning reflection + planning	Reset goals for the week
10:00–1:00 PM	Study catch-up or project work	Stay ahead of deadlines
1:00–2:00 PM	Lunch + break	Rest
2:00–5:00 PM	Personal development (watch lectures, portfolio, or AI tools)	Build future-ready skills
Evening	Social/family time or hobby	Reconnect & relax

Sunday

Time	Activity	Purpose
Morning	Light chores, journaling, spiritual/mindful practice	Refresh for a new week
Afternoon	Outdoor activity or complete pending tasks	Balance and productivity
Evening	Plan upcoming week & early bedtime	Prepare confidently for Monday

Tips for Success

- Use the Pomodoro Technique (25-min focus + 5-min break) during study blocks.
- Keep Sundays screen-light to prevent burnout.
- Remember: **Consistency > Intensity**. Even small steps daily create big wins over time.
- Celebrate progress—*every assignment completed is a step toward your dream career!*

Figure 2. Test 1 – Default Chat Output (Weekly Time-Management Plan, Part 2)

Test 2 - Friendly Tone Chat

Modified Prompt:

“You are an encouraging AI coach helping a busy university student balance study, part-time work, and personal wellness. Create a one-week schedule that feels supportive, motivational, and realistic.”

Output Summary:

The AI created a **“Growth & Gratitude”-themed weekly plan** filled with thoughtful affirmations and reflection points. Each day included not only structured time blocks but also positive reminders like *“Gratitude fuels growth”* and *“Each week is a new chance to shine brighter.”* The model added balanced routines that considered both productivity and emotional wellness — such as morning gratitude reflections, evening journaling, and mindfulness breaks.

Observation:

This version was **emotionally engaging and relatable**, with a coaching style that encouraged motivation and positivity. The addition of mantras, quotes, and rest-focused messages showed how AI can enhance user experience through empathy. It effectively humanized the schedule, though it leaned more on motivation than task optimization.

Assessment: 9 / 10

Screenshot:




Figure 3. Test 2 – Friendly-Tone Chat Output (Weekly Focus Schedule – Part 1)




Figure 4. Test 2 – Friendly-Tone Chat Output (Weekly Focus Schedule – Part 2)

Test 3 – Professional Focus Chat

Modified Prompt:

“You are a productivity-focused AI assistant helping a university student manage time effectively while taking online courses and working part-time. Create a one-week schedule that uses productivity techniques such as the Pomodoro Method and SMART goals to increase focus and efficiency. Keep the tone professional and structured, and include specific daily objectives, short breaks, and clear milestones for tracking progress.”

Output Summary:

The LLM produced a highly structured **Weekly Productivity Schedule** using **SMART goals** and the **Pomodoro Method**. The plan began with an explicit weekly objective — maximizing efficiency and achieving measurable academic progress — followed by a SMART goal breakdown (Specific, Measurable, Achievable, Relevant, Time-bound). It included a detailed **day-by-day table** outlining study sessions, Pomodoro cycles, and milestones such as discussion post completion or assignment submission.

The assistant also provided a **Progress Tracking Template** and **Pro Tips**, recommending productivity apps like *Forest* and *Focus To-Do* and introducing reflective routines for self-assessment.

Observation:

This version felt **professional, technical, and action-driven**. Unlike the friendly tone of Test 2, this response focused on measurable outcomes, clear milestones, and time efficiency. The structured format, bullet points, and quantifiable metrics gave it a polished project management feel — ideal for productivity-oriented users. However, it lacked emotional warmth, which made it less motivational but more business-like.

Assessment: 8 / 10

Screenshot:

You are a productivity-focused AI assistant helping a university student manage time effectively while taking online courses and working part-time.

Create a one-week schedule that uses productivity techniques such as the Pomodoro Method and SMART goals to increase focus and efficiency.

Keep the tone professional and structured, and include specific daily objectives, short breaks, and clear milestones for tracking progress.

Here's a **structured one-week productivity schedule** tailored for a busy university student balancing online coursework and part-time work, incorporating the **Pomodoro Technique**, **SMART goals**, and clear milestones.

 **Weekly Productivity Schedule**

Objective:
Maximize study efficiency, maintain work-life balance, and achieve measurable academic progress through time-blocking and focus strategies.

SMART Goal for the Week

Specific: Complete all assigned readings, two major coursework submissions, and one reflection post.
Measurable: Minimum of 12 study sessions (25-minute Pomodoro intervals × 4 per day).
Achievable: Spread study blocks across five days.
Relevant: Directly aligned with weekly academic deliverables.
Time-bound: Complete by Sunday 8:00 PM.

Daily Framework

- **Pomodoro Cycle:** 25 minutes focused work + 5-minute short break
- **Extended Break:** 15–20 minutes after every 4 Pomodoros
- **Study Milestone Tracker:** ☒ Each Pomodoro session completed = +1 milestone

Figure 5. Test 3 – Professional Focus Chat Output (Weekly Productivity Schedule Part- 1)



Weekly Schedule Overview

Day	Focus Areas & Tasks	Study Blocks (Pomodoros)	Milestones / Notes	
Monday	- Review lecture notes - Plan the week using a task list - Begin readings for Module 1	4 Pomodoros	✔ Task list finalized ✔ Notes summarized	
Tuesday	- Complete discussion post draft - Watch video lectures	4 Pomodoros	✔ Discussion post ready for review	
Wednesday	- Revise Tuesday's draft - Practice quiz or mini test	3 Pomodoros	✔ 1 quiz completed	
Thursday	- Begin major assignment (Part 1) - Outline key arguments	4 Pomodoros	✔ Assignment outline submitted for feedback	
Friday	- Continue assignment (Part 2) - Edit and format report	4 Pomodoros	✔ First draft completed	
Saturday	- Proofread final submission - Submit coursework - Light review of next module	3 Pomodoros	✔ Submission done ✔ Next module previewed	
Sunday	- Reflection post - Plan next week's goals - Digital cleanup	2 Pomodoros	✔ Week reflection completed	

Break Routine Suggestions

- **Short Breaks (5 min):** Stretch, hydrate, or walk around your room.
- **Long Breaks (15–20 min):** Quick meal, light exercise, or mindfulness breathing.
- **Evening Wind-Down (30–45 min):** Journaling or reading to reset focus.

Figure 6. Test 3 – Professional Focus Chat Output (Weekly Productivity Schedule Part- 2)

Progress Tracking Template

Goal	Target	Achieved	Status	
Study sessions completed	20 Pomodoros		☐	
Discussion post submitted	Tuesday		☐	
Major assignment submitted	Saturday		☐	
Reflection post submitted	Sunday		☐	

Pro Tips

- Use apps like **Forest** or **Focus To-Do** to track Pomodoros.
- Keep a **"Distraction Log"** to note interruptions and improve focus.
- End each day with a **2-minute reflection**: "What worked well? What needs adjusting?"

Figure 7. Test 3 – Professional Focus Chat Output (Weekly Productivity Schedule Part- 3)

Overall Reflection – Section 1: Practice with an LLM

Running the three tests in separate chats helped me clearly see how **prompt design and tone** shape the quality of an AI's responses. Although each prompt employed the same core task, the results were strikingly different. The **default chat** produced a structured but generic plan; the **friendly-tone chat** added empathy, motivation, and human warmth; and the **professional-focus chat** delivered measurable goals, detailed routines, and strong time-management logic.

This comparison showed that large language models are highly sensitive to wording and intent. Small adjustments, such as asking for encouragement versus structure, can completely change the model's personality and output style. I also realized that there is no single "best" version: the right tone depends on the user's needs. Overall, this activity improved my ability to craft precise prompts and evaluate AI output based on clarity, usefulness, and emotional impact.

Feature	Test 1 (Default)	Test 2 (Friendly)	Test 3 (Professional)
Tone	Neutral	Supportive	Structured
Focus	General balance	Motivation & empathy	Efficiency & metrics
Personalization	Low	High	Moderate
Creativity	Moderate	High	Moderate
Practical Application	Good	Good	Excellent

Section 2 – Practice with a Custom GPT

Tool Used

Custom GPT: Expedia GPT (Custom GPT by Expedia Group)

Context:

For this exercise, I chose to interact with **Expedia GPT**, a travel-planning assistant designed to recommend destinations, hotels, and itineraries based on user preferences and budgets. I wanted to explore how a domain-specific GPT applies real-world data and reasoning to build personalized travel plans.

I started by providing the following prompt:

“Hi Expedia! I’m planning a 6-day trip from Dallas, Texas, to San Francisco, California, next month, after November 20th, anytime for multiple tech-based conferences. I’d like help creating a balanced itinerary that mixes sightseeing, food, and photography. Please include approximate budget and travel suggestions.”

The goal was to test how effectively the GPT could combine conference-related travel planning with personal interests, such as sightseeing and photography, while staying within a practical budget.




Figure 8. Initial Interaction with Expedia GPT – Trip Planning Prompt and Response

Observation:

The response was professional and well-organized. The GPT clearly demonstrated an understanding of my travel priorities, combining business and leisure effectively. It also displayed adaptive questioning by asking for clarification on neighborhood preferences before proceeding, like how a human travel consultant would interact. The tone felt polite, knowledgeable, and action-driven, which set a strong base for deeper interaction.

Interaction 2 – Adjusting for Affordability**Follow-Up Prompt:**

Can you make the plan more affordable by suggesting budget-friendly hotels? I'm flexible about where I stay since I plan to rent a car, so it'll be easy to move around.

Can you make the plan more affordable by suggesting budget-friendly hotels? I'm flexible about where I stay since I plan to rent a car, so it'll be easy to move around.

Talked to apim.expedia.com ▾

Here are three **budget-friendly lodging options in San Francisco** I've found for your 6-day stay (Nov 24 – 30, 2025).

I assumed you prefer **affordable hotels or hostels** within easy reach of the city's main attractions, and that you're flexible about location since you'll have a rental car. (We can adjust the search if you'd like to focus on a specific area or upgrade to mid-range later!)

 **1. European Hostel**



A basic, backpacker-style hostel near Bill Graham Civic Auditorium — ideal if you want to save money and stay central.

- ★ **Star Rating:** 1.5
- 🗨 **Guest Rating:** 5.4 / 10 (660 reviews)
- 🛏 **Max Occupancy:** 2
- 📍 **Location:** Downtown San Francisco
- 💰 **Average Nightly Price:** \$42.75
- 👛 **Total (6 nights, taxes & fees included):** \$301.44 USD

Figure 9. Expedia GPT Recommending Budget-Friendly Lodging Options

Output Summary:

The GPT connected to Expedia's API and returned **three detailed, low-cost lodging options** for November 24–30, 2025. Each option included a star rating, guest review score, price breakdown, and location summary:

1. **European Hostel** – Central location, lowest price at **\$301.44 total**.
2. **Chapter San Francisco** – Highly rated Victorian-style hostel with modern amenities (**\$365.80 total**).
3. **Hotel North Beach** – Near the Embarcadero and Chinatown (**\$389.10 total**).

It emphasized that each choice kept the total budget under \$400, freeing up funds for flights, food, and experiences. The assistant also acknowledged my flexibility with location and offered to continue by finding flight options.

Observation:

The GPT adapted perfectly to my new constraints. It shifted from general travel planning to **cost-focused reasoning**, providing transparent price data and linking its choices to my stated needs (flexibility and photography). The tone remained professional yet conversational. What stood out was its ability to explain the **financial impact** of my change in preference — a level of awareness that made the exchange feel practical and personalized.

Interaction 3 – Asking for More or Better

Follow-Up Prompt:

Can you also recommend a few must-see spots for photography in San Francisco — both classic and hidden gems?




Figure 10. Expedia GPT Suggesting Classic and Hidden Photography Spots in San Francisco

Output**Summary:**

The GPT produced a curated list of **San Francisco photography locations**, neatly divided into two categories:

- **Classic Spots:** Included landmarks such as the Golden Gate Bridge, Alamo Square, Pier 7, Twin Peaks, and the Palace of Fine Arts — each with specific photography advice, *best shooting times*, and composition tips.
- **Hidden Gems:** Featured less-crowded areas like Lands End, the 16th Avenue Tiled Steps, Clarion Alley murals, Bernal Heights Swing, Fort Point, and Crissy Field.

The output balanced tourist favorites with local recommendations, catering to both sightseeing and creative photography interests. It even mentioned **lighting conditions**, **mood composition**, and **drone-friendly zones**, showing contextual awareness beyond basic travel suggestions.

Observation:

This interaction highlighted how well the GPT could expand and **enrich a travel plan** when asked for more detail. Instead of just listing places, it explained why and when each location works best for photographers — combining **technical guidance** (e.g., “best time: golden hour”) with **local insights**. The assistant’s tone remained friendly and informative, showing excellent adaptability from a budget conversation to a more creative and personalized one.

Interaction 4 – Understanding the GPT’s Reasoning

Prompt Used:

“Why did you recommend these specific photography locations over others? What criteria did you use?”

Output Summary:

The GPT responded with a structured, five-point explanation describing its selection process for San Francisco photo spots.

It stated that the choices were based on:

1. **Visual Variety & Composition Opportunities** – balancing architecture, coastal views, and color tones to enrich a travel photo portfolio.
2. **Lighting & Atmosphere Potential** – considering time of day and lighting direction, highlighting locations ideal for sunrise, sunset, or blue hour photography.
3. **Accessibility & Practicality** – including both central and drive-to locations, taking into account my note about renting a car.
4. **Local & Hidden-Gem Factor** – ensuring authentic, less-crowded experiences that capture San Francisco’s local spirit.
5. **Storytelling Potential** – emphasizing sites that reflect the city’s cultural blend of technology, art, and natural beauty.

Observation:

This interaction clearly showed that **Expedia GPT is capable of transparent reasoning**, breaking down how its recommendations align with user context. Its structured explanation felt similar to how a human travel planner justifies choices—mixing logic with creativity. The reference to my rental car demonstrated context retention and adaptability, while the storytelling component revealed a sophisticated awareness of user intent (photography, culture, and exploration).

Screenshot:

Why did you recommend these specific photography locations over others? What criteria did you use?

That's a great question — and I love that you're thinking like a photographer 📷

Here's exactly **how I chose those San Francisco photo locations** and the **criteria** behind the mix of classics and hidden gems:

📍 1. Visual Variety & Composition Opportunities

I looked for places that offer **diverse visual elements** — architecture, water reflections, skyline views, color, and texture — so you can build a rich photo portfolio during your trip.

- **Examples:**
 - *Golden Gate Bridge* → iconic symmetry and natural fog effects.
 - *Clarion Alley* → bold colors and urban storytelling.
 - *Lands End* → dramatic coastline for long exposures.

🕒 2. Lighting & Atmosphere Potential

Every spot was selected with **light direction and timing** in mind — sunrise/sunset access, golden hour reflections, and low-light city scenes.

- *Twin Peaks* and *Battery Spencer* are unbeatable for sunrise/sunset.
- *Pier 7* and *Palace of Fine Arts* shine at blue hour for long exposure work.

🚶 3. Accessibility & Practicality

Because you mentioned you'll have a **rental car**, I included both easy-to-reach city spots and short drives that make the most of your mobility (like *Muir Woods* and *Half Moon Bay*).

If you were relying on public transit, I'd focus more tightly around SoMa, the Embarcadero, and Chinatown.

🕒 4. Local & Hidden-Gem Factor

Beyond the famous icons, I wanted you to experience **the real, less-touristy San Francisco** that locals love photographing:

- *Bernal Heights Swing* → intimate skyline view with few crowds.
- *16th Avenue Tiled Steps* → colorful pattern compositions, often missed by first-time visitors.

🎨 5. Storytelling Potential

Each spot helps tell a piece of San Francisco's identity — its tech-meets-art culture, Victorian charm, and Pacific drama.

The combination creates a **photo narrative** that transitions from cityscape to nature to local culture, ideal for a travel or photography blog, or even for conference downtime exploration.

Figure 11. Expedia GPT Explaining Its Selection Criteria for Photography Recommendations

Interaction 5 – Flight Recommendations

Prompt Used:

“Can you recommend affordable round-trip flight options from Dallas (DFW or DAL) to San Francisco (SFO or OAK) for Nov 24–30, 2025? Include prices, times, and airlines.”

Output Summary:

Expedia GPT noted that the API supports **one-way** searches only, then returned **outbound** options for **Nov 24, 2025** with clear reasoning (lowest price, non-stop preference, convenience) and suggested pairing with a return on Nov 30:

1. **Frontier 3999** – DFW→SFO 12:06–2:17 pm, **4h11m non-stop**, **\$127.98**
2. **United 1184** – DFW→SFO 7:49–9:50 pm, **4h01m non-stop**, **\$129.49**
3. **Southwest (via SAN)** – DAL→SFO 6:30 am–12:10 pm, **7h40m 1-stop**, **\$139.18** (bags included)

It also added a tip that **round-trip on Expedia.com often saves ~10–15%** versus two one-ways, and offered to fetch **return** options next.

Observation:

Great **constraint handling**: when round-trip wasn’t available via API, it shifted to one-way with a clear pairing strategy and budget guidance. The mix of non-stop vs 1-stop and airline trade-offs showed practical reasoning aligned to my cost/time goals.




Figure 12. Expedia GPT Generating Affordable Flight Options from Dallas to San Francisco

Interaction 6 – “What if...?” Scenario (do this next)

Prompt Used:

“What if I fly into Oakland (OAK) instead of San Francisco (SFO) to save money? Can you show me flight options for that route?”

Output Summary:

The GPT returned a full list of **one-way flights from Dallas Love Field (DAL) to Oakland (OAK)** for **Nov 24, 2025**, explaining that round-trip searches require Expedia.com.

It suggested three affordable options—all **Southwest Airlines**—with clear trade-offs:

1. **DAL → OAK (Non-stop, 8:55 PM – 10:45 PM)** 3 h 50 m at **\$163.38 USD**
2. **DAL → OAK (via Burbank)** 6 h 05 m at **\$216.64 USD**
3. **DAL → OAK (via Las Vegas)** 5 h 25 m at **\$226.77 USD**

The assistant added a **ground-transport comparison**, outlining:

- **BART train:** ≈ \$10 one way (45–50 min)
- **Rental car:** \$40–70/day near OAK
- **Rideshare/taxi:** \$45–60 to downtown SF

It even summarized that OAK flights might cost slightly more than SFO in airfare but could **save time and reduce airport congestion**, showing strong analytical reasoning.

Observation:

This exchange demonstrated the GPT’s ability to **adapt plans dynamically** and provide **data-driven trade-offs** when given a “what-if” condition. It retained trip context (rental-car use, travel dates, and goals) and produced comparative insights—balancing time, cost, and convenience. The assistant’s tone stayed helpful and concise, mirroring a professional travel consultant’s explanation of pros and cons.

Screenshot:




Figure 13. Expedia GPT Adapting Flight and Ground Transport Options for an Alternate Airport (OAK)

Interaction 7 – Conversation Summary

Prompt Used:

“Please summarize our full trip conversation into a clear 6-day plan with lodging, flights, photography highlights, estimated total cost range, and travel tips.”

Output Summary:

The GPT delivered a complete **6-day Dallas-to-San Francisco itinerary**, integrating every earlier element:

- **Flights:** Frontier, United, and Southwest options (\$260 – \$320 round-trip) with timing and airport comparisons (SFO vs OAK).
- **Lodging:** Three affordable choices (European Hostel \$301, Chapter SF \$365, Hotel North Beach \$389).
- **Photography Plan:** Day-by-day highlights from Pier 7 and Golden Gate Bridge to Clarion Alley murals and Lands End sunsets.
- **Transport:** Rental-car and public-transit cost breakdowns.
- **Budget Summary:** Estimated total \$960 – \$1,270 per person for 6 days, with travel tips on weather, parking, timing, and gear.

Observation:

This final summary demonstrated **strong contextual memory and synthesis**. The GPT accurately merged logistics, creative elements, and costs into one cohesive plan, maintaining the same friendly-consultant tone seen throughout the session. Its ability to recall prior constraints (budget, car rental, photography focus) and express them in a clear, human-readable itinerary showed excellent comprehension and adaptability.

Overall Reflection – Section 2: Practice with a Custom GPT

Working with **Expedia GPT** felt like collaborating with a professional travel planner powered by AI. Each new question — from budget adjustments to alternate airports — was met with clear logic, data, and actionable insights. I learned how **domain-specific GPTs combine structured APIs with conversational reasoning** to create personalized, data-driven plans. The model’s transparency about its limits (one-way flight data) and its ability to adapt when I changed goals reflected practical problem-solving in real time. This exercise deepened my understanding of how custom GPTs apply generative AI to specialized industries, blending **automation, empathy, and user-context awareness** to enhance decision-making and user experience.

Please summarize our full trip conversation into a clear 6-day plan with lodging, flights, photography highlights, estimated total cost range, and travel tips.

Here's your **6-day San Francisco itinerary summary** — a complete and affordable plan blending your tech conference goals with sightseeing, food, and photography adventures 🇺🇸 📷

📅 Trip Overview

Route: Dallas, TX → San Francisco, CA

Dates: November 24 – 30, 2025 (6 days / 5 nights)

Focus: Tech conferences, city exploration, food, and photography

✈️ Flights (Dallas → Bay Area)

Option 1: San Francisco International (SFO)

- **Frontier Airlines:** Non-stop, departs DFW 12:06 PM, arrives SFO 2:17 PM
💰 \$127.98 USD (one-way)
 - **United Airlines:** Non-stop, departs DFW 7:49 PM, arrives SFO 9:50 PM
💰 \$129.49 USD (one-way)
- 🔗 [View flights on Expedia ↗](#)

Option 2: Oakland International (OAK)

(Slightly higher fare but often faster customs and less traffic)

- **Southwest Airlines:** Non-stop from DAL 8:55 PM → OAK 10:45 PM
💰 \$163.38 USD (one-way)
- 🔗 [View flights on Expedia ↗](#)

📊 **Estimated Round-Trip Total:** \$260 – \$320 USD per person (including taxes & fees)

🏠 Lodging (Nov 24–30, 2025)

All taxes & fees included in total price quotes.

Hotel	Area	Type	Guest Rating	Total (6 nights)
European Hostel	Downtown	Hostel	5.4 / 10	💰 \$301.44 USD
Chapter San Francisco	Civic Center	Boutique Hostel	9.2 / 10	💰 \$365.80 USD
Hotel North Beach	Embarcadero	Budget Hotel	7.0 / 10	💰 \$389.10 USD

🔗 [View hotels on Expedia ↗](#)

📊 **Average Lodging Budget:** \$320 – \$400 USD total per person

Figure 14. Expedia GPT Compiling the Final 6-Day San Francisco Itinerary

Section 3 - Practice with a Research Assistant (STORM AI)

Tool Used:

STORM AI (Stanford OVAL Lab)

Topic: *The Evolution of Big Data Analytics*

Context:

For this activity, I utilized STORM AI, an open-source research assistant system developed at Stanford University's Open Virtual Assistant Lab. The tool is designed to synthesize information from multiple scholarly sources and generate well-cited research articles. I selected the topic "*The Evolution of Big Data Analytics*" because I have previously written on this subject and could evaluate the AI's accuracy and depth against my own understanding.

Interaction and Output Summary

Prompt Used:

"Generate a research article titled *The Evolution of Big Data Analytics*, covering historical development, key frameworks, and modern trends.

Generate a research article titled The Evolution of Big Data Analytics, covering historical development, key frameworks, and modern trends.

Generating your article (it may take 3 minutes)•••

Finish browsing <https://rikkeisoft.com/blog/data-analytics-framework/>
 Finish browsing <https://www.ibm.com/think/topics/big-data-analytics>
 Finish browsing <https://medium.com/technology-nineleaps/the-evolution-of-data-analytic-s-539ea532a90d>
 Finish browsing <https://fiveable.me/lists/key-concepts-in-big-data-processing-frameworks>
 Finish browsing <https://www.bigdataframework.org/short-history-of-big-data/>
 Finish browsing <https://www.3pillarglobal.com/insights/blog/the-evolution-of-big-data-analytics-market/>
 Finish browsing <https://trainingcred.com/blog/big-data-technologies-and-frameworks-the-future-of-scalable-data-processing>
 Finish browsing <https://www.geeksforgeeks.org/blogs/big-data-frameworks/>
 Finish browsing <https://motherduck.com/learn-more/big-data/>
 Finish browsing <https://jelvix.com/blog/top-5-big-data-frameworks>
 Finish browsing <https://utd-ir.tdl.org/bitstreams/ec2ca5dd-3325-4722-8f1f-0cd7610c52c3/download>
 Finish browsing <https://medium.com/womenintechology/big-data-architectures-key-frameworks-for-scalable-data-processing-fa7948ab8922>
 Finish browsing <https://data.folio3.com/blog/data-analytics-framework/>
 Finish browsing <https://www.turing.com/resources/best-big-data-platforms>
 Finish browsing <https://www.oracle.com/a/ocom/docs/big-data/big-data-evolution.pdf>
 Finish browsing <https://www.lightsondata.com/the-history-of-big-data/>

Figure 15. STORM AI Generating a Research Article on "*The Evolution of Big Data Analytics*"

Data Scientist Historian of Technology Basic fact writer Business Analyst

This editor will concentrate on key frameworks and methodologies used in big data analytics today. They will explain statistical techniques, machine learning algorithms, and data processing frameworks like Hadoop and Spark, illustrating how these tools have evolved over time.



What were the key milestones in the historical development of big data analytics that shaped its current state?



The historical development of big data analytics is marked by several key milestones that have significantly shaped its current state. 1. **Early Statistical Analysis**: The history of data analysis can be traced back as far as the 1600s, when statistical analysis began to emerge as a tool for understanding data and solving real-life problems. This foundational work laid the groundwork for future data-driven approaches (Sources [10], [39]). 2. **The First Programmable Computers**: The introduction of the first programmable computer in the 1940s marked a significant leap in data processing capabilities. This technology was initially used for complex calculations, such as deciphering Nazi codes during World War II, demonstrating the potential of electronic computing for data analysis (Sources [10], [14]). 3. **Emergence of Data Processing (1950s-1960s)**: The 1950s and 1960s saw the rise of data processing technologies, which included the development of high-speed calculations and the first electronic data processing machines. This era set the stage for more sophisticated data management techniques (Sources [13], [18]). 4. **Business Intelligence Rise (1970s-1980s)**: During the 1970s and 1980s, the concept of business intelligence began to take shape, focusing on data analysis to inform business decision-making. This period marked the transition from mere data processing to strategic data usage (Sources [13], [31]). 5. **Introduction of Big Data Concepts (1990s)**: The term "big data" began gaining traction in the 1990s as businesses started recognizing the limitations of traditional data management systems in handling large volumes of information. This decade witnessed the rise of parallel database management systems, catering specifically to big data needs (Sources [10], [23]). 6. **Advancement of Big Data Technologies (2000s)**: The 2000s saw significant advancements in big data technologies, including the emergence of Hadoop, which provided an open-source framework for processing large datasets. This technology revolutionized data storage and processing, allowing organizations to analyze vast amounts of data efficiently (Sources [10], [36]). 7. **Big Data 3.0 and Modern Trends (2010s-Present)**: The current era, referred to as Big Data 3.0, began around 2010 with the advent of smartphones, GPS sensors, and the cloud.




Can you elaborate on the key frameworks and methodologies currently used in big data analytics, specifically mentioning any notable statistical techniques and machine learning algorithms?



In the realm of big data analytics, various key frameworks and methodologies have emerged

Figure 16. STORM AI Generating and Explaining Research Insights on Big Data Analytics

This image shows STORM AI in action as it synthesizes historical milestones, frameworks, and machine-learning methodologies in Big Data Analytics. The tool adopts multiple expert personas, Data Scientist, Historian of Technology, Fact Writer, and Business Analyst, to generate multi-layered responses. This illustrates STORM AI's ability to reason across perspectives, cite sources, and provide structured, academically relevant content.



The screenshot displays the STORM AI interface. On the left is a sidebar with navigation options: '+ New Session', 'Discover', 'My Library', and 'Table of Contents'. The 'Table of Contents' is expanded, showing a 'summary' section with 'Historical Development' (Early Foundations of Data Collection, Development of Statistical Methods, Technological Innovations in the 19th and 20th Centuries, The Rise of Big Data) and 'Key Frameworks' (Overview of Data Analytics Frameworks, Types of Data Analytics Frameworks, Distributed Computing Frameworks, Apache Hadoop, Apache Spark, Emerging Trends in Data Analytics Frameworks). Below these are 'Modern Trends' (The Proliferation). The main content area is titled 'STORM' and contains a detailed article on data analytics frameworks. The article starts with an introduction, followed by sections on 'Data Analytics Strategy Framework', 'Data Analytics Governance Framework', and 'Data Analytics Capability Framework'. It then discusses 'Distributed Computing Frameworks', 'Apache Hadoop', and 'Apache Spark'. The article is well-structured with clear headings and subheadings, and includes a 'References' section at the bottom.

Figure 17. Structured Research Output from STORM AI

Output Summary:

STORM AI produced a detailed, 10-section article of about 2,000 words. It began with a comprehensive **historical overview**, tracing data collection from the Sumerians' clay tablet records to the formalization of statistics in the 17th century and the mechanization of data processing in the 19th and 20th centuries.

The article then covered the **rise of Big Data** in the late 20th century, identifying the exponential growth of digital information as the key driver for scalable analytics. It outlined several **core frameworks, including Data Analytics Strategy, Governance, and Capability Frameworks**, before introducing distributed-computing systems such as **Apache Hadoop** and **Apache Spark**, explaining their architecture and use cases.

The final sections discussed **modern trends**, including cloud computing, AI, and machine-learning integration, mobile data proliferation, democratization of analytics, and emerging technologies like blockchain, edge computing, and quantum analytics. It also examined **challenges and ethics**, including data quality, privacy, cost, and cybersecurity. The paper concluded with **future directions**, emphasizing the importance of prescriptive analytics, data democratization, and the ethical use of AI in data analysis.

How I Think STORM AI Works

From observing the brainstorming trace, STORM AI appears to use a **retrieval-and-synthesis pipeline**. It breaks the user prompt into sub-questions (e.g., “*What are the major Big Data frameworks?*”, “*When did Hadoop emerge?*”) and retrieves information from multiple indexed sources. Using its **Vector-RM** retrieval module, it grounds statements in related documents, then composes a draft through iterative summarization. The model applies a citation-insertion step to link each factual segment to supporting references. This process reflects **multi-perspective question-asking**, allowing it to merge historical, technical, and ethical dimensions into a coherent narrative.

Observation

The generated article was coherent, chronological, and technically sound. It accurately highlighted key eras (Hadoop → Spark → Cloud → AI Integration) and used domain-appropriate terminology. I compared it with my own earlier article and found a strong overlap in structure, but STORM's version offered broader citation coverage and more emphasis on ethics and future research. A few sections read slightly mechanically, but overall, it demonstrated impressive contextual reasoning and evidence-based synthesis.

My Article on the same topic: <https://medium.com/@sandeep-sudesh-kumar/evolution-of-big-data-analytics-25daf091ffad>

Reflection

Using STORM AI revealed how research-oriented AI tools differ from conversational models. Instead of producing quick summaries, STORM performs evidence-driven composition, creating academic-style writing with verifiable citations. This experience showed me how such systems can accelerate literature reviews, enhance academic rigor, and support collaboration between human insight and AI efficiency.

I compared the AI-generated article with my own previously published piece titled “*The Evolution of Big Data Analytics*” (authored under my name). While my version emphasized the business implications and real-world adoption of Big Data, STORM AI's version focused more on academic structure, chronological progression, and citation variety. It incorporated greater historical depth, covering early statistical developments and computing milestones, which made it feel more like a formal research review. This comparison helped me recognize how human authors bring narrative and contextual perspective, while AI tools excel at organization, breadth, and speed of synthesis.

Overall, I realized that while AI can efficiently gather and structure knowledge, the human element remains vital for interpretation, ethical reflection, and storytelling that connects data to meaning.

Section 4 – Design Thinking and Empathy in AI Problem-Solving

Context

This section explores how **design thinking** and **empathy** help guide AI/ML problem-solving to create ethical, user-centered, and effective solutions. Design thinking involves five key stages: **Empathize, Define, Ideate, Prototype, and Test**, which together ensure that technological innovation aligns with real human needs and fairness principles.

My Approach to Problem-Solving

When faced with a problem, my first step is to **understand it thoroughly before rushing to a solution**. I take time to break down the issue into smaller components, identify possible causes, and analyze who is affected and how.

Rather than focusing immediately on outcomes, I begin by clarifying *why* the problem exists. Once I understand the “why,” I brainstorm several possible solutions, compare them based on feasibility, and select the one most aligned with the goal.

To determine whether a solution is working, I rely on feedback loops, measurable results tracking data, observing user responses, and continuously refining the approach until it achieves consistency and a positive impact.

User-Centric Thinking

I believe that effective AI systems are designed for people first. Whenever I approach a project, I ensure that I consider the users’ perspectives, limitations, and goals.

In my experience, some technical solutions succeed functionally but fail in practice because they don’t meet real-world user expectations, for example, dashboards that provide accurate data but are too complex for end users.

To prevent this, I prioritize empathy-driven approaches such as user interviews, observation, and feedback testing. These help me connect emotionally and intellectually with the people experiencing the problem, ensuring the solution reflects their lived reality.

Understanding Empathy in AI/ML

In the context of AI and machine learning, empathy ensures technology serves people fairly and transparently. Empathy in AI means anticipating how algorithms impact users — especially in sensitive areas like bias, privacy, and accessibility. Starting with empathy helps define problems accurately before developing models or solutions.

By identifying users’ emotions, frustrations, and needs early, I can avoid ethical pitfalls such as biased datasets or unfair decision-making. This ensures that AI systems enhance trust and inclusion rather than amplifying inequities.

Applying Design Thinking to AI

The Design Thinking Process provides a structured yet flexible framework for solving complex AI problems.

Phase	Focus	Example in AI/ML Context
Empathize	Understand user pain points through interviews, observation, or surveys.	Interview data scientists and business users to understand challenges in interpreting AI model outputs.
Define	Frame the problem based on user insights and feasibility.	Translate user frustration (“I can’t understand the model results”) into a clear problem statement: “Lack of explainability in AI decisions.”
Ideate	Brainstorm multiple AI-powered solutions.	Propose solutions like visualization dashboards, natural-language explanations, or bias-detection models.
Prototype	Build quick models or mock-ups to test.	Create a small interface or sample dataset to test the explainability dashboard.
Test	Validate with real users; gather feedback.	Test the prototype with stakeholders, check usability, fairness, and clarity, and iterate to improve it.

Reflection

Working through this section helped me see how empathy bridges technology and humanity. In AI/ML, it’s easy to become absorbed in algorithms and data accuracy, but design thinking reminds us that behind every dataset are people whose lives are impacted by these systems.

By integrating empathy, feedback, and ethical reflection throughout the design process, I can ensure that any AI solution I create is not just functional but responsible, fair, and meaningful.

Section 5 – Building an AI Chatbot or Assistant

Context

For this lab, I designed and built a custom AI productivity assistant named FocusMate, using the Design Thinking process. The objective was to create a friendly, motivational chatbot that helps students and professionals manage their study or work schedules, stay focused, and maintain consistency through personalized productivity frameworks such as the Pomodoro Technique and SMART goals.

FocusMate was built in ChatGPT's Custom GPT builder, with the aim of blending empathetic dialogue, time-management strategies, and mental-wellness support into a single, conversational AI tool.

Planning the Bot – Using Design Thinking

1. What is your bot's objective?

To create a Personal Productivity & Motivation Assistant called FocusMate that helps students and professionals plan their day, stay organized, and stay motivated through supportive, empathy-driven conversations.

2. What kind of information should the assistant use?

- Time-management frameworks (Pomodoro Technique, SMART Goals)
- Motivational reminders and reflection prompts
- Wellness suggestions (short breaks, breathing, stretching exercises)

3. What should the assistant *not* do?

- Provide medical, psychological, or financial advice
- Store, recall, or process personal data
- Use negative, harsh, or judgmental language

4. How should the assistant engage with the user?

By maintaining a friendly, encouraging, and conversational tone, asking clarifying questions, providing structured plans, and ending each exchange with a short motivational or affirming message.

5. Will images or other data be part of the interaction?

No images or media will be used in the prototype version.
(Future iterations could display productivity dashboards, visual progress trackers, or charts for user reflection.)

6. How should the assistant start?

It begins with a warm, positive greeting that encourages planning:

“Hey there! What’s your main goal or task for today? Let’s plan how to tackle it together.”

7. Are there steps the bot should follow?

1. Greet the user and show empathy
2. Identify the user’s main goal or challenge
3. Suggest a plan using Pomodoro or SMART goal principles
4. Offer encouragement and short breaks
5. Summarize key next steps

8. How should the assistant reach a conclusion?

By providing a brief summary of the day’s plan and a positive affirmation, such as:

“Perfect — two Pomodoro sessions before lunch. You’ve got this! Stay focused and reward yourself after!”

9. What does success look like?

A successful interaction is when the user:

- Feels calm, focused, and less overwhelmed
- Has a clear, realistic action plan for the day
- Leaves the conversation feeling supported and motivated

10. Scenario to Test

User Input:

“I have three assignments due and can’t focus.”

Expected Bot Behavior:

1. Acknowledge the user’s stress with empathy
2. Break tasks into smaller, timed sessions (Pomodoro format)
3. Suggest short breaks between sessions
4. End with motivation and reassurance:

“Let’s turn that stress into structure. You’re doing great - just one step at a time!”

Build Tool Selection

Tool Used: ChatGPT (Custom GPT Builder)

Privacy Review:

- Chats can be deleted; data is not publicly stored.
- Clear terms for educational use and private projects.

Rationale:

ChatGPT was chosen for its ease of setup, flexible tone control, and ability to simulate realistic AI assistant behavior without requiring programming knowledge.

Execution and Testing

Process

1. Built FocusMate with a coffee-cup icon to reflect warmth and productivity.
2. Configured behavior rules emphasizing empathy, brevity, and structured goal support.
3. Tested three core interaction types and refined for shorter, natural responses.

Step-by-Step: Building the “FocusMate” AI Assistant

Step 1: Open the GPT Builder

1. Visit <https://chat.openai.com/gpts>
2. Click “Explore GPTs” → “Create”
3. The GPT Builder interface includes:
 - Create Tab: For natural-language setup
 - Configure Tab: For technical settings

Step 2: Give Your Bot an Identity

In the *Create* tab, type:

“Build a friendly productivity and motivation coach called FocusMate. It helps students and professionals plan their day, stay organized, and feel motivated using methods like Pomodoro and SMART goals.”

The builder automatically generates a short description and icon.




Figure 18. Building FocusMate in the Custom GPT Builder — showing the creation phase where the chatbot's name, personality, and profile concept ("a coffee cup with a to-do list") were defined to represent warmth, focus, and motivation.

Step 3: Provide Behavioral Instructions

Set tone and behavior rules:

- Always start with empathy.
- Keep responses under 5 sentences unless detailed planning is required.
- Use motivational and inclusive language.
- Avoid over-explaining, guide with calm clarity.




Figure 19. Configuring FocusMate's behavior rules — defining empathy, structure, and motivational tone. The setup ensures the chatbot starts with encouragement, asks clarifying questions, uses frameworks like Pomodoro or SMART goals, and avoids negative comments.

Step 4: Add Conversation Starters

In Configure → Conversation Starters, include examples such as:

- “Help me organize my study schedule for this week.”
- “I’m overwhelmed by deadlines. Can you make a plan?”
- “Remind me to take focus breaks using Pomodoro.”
- “Give me a quick pep talk before I start studying.”

These act as intuitive entry points for new users.




Figure 20. Final configuration of FocusMate in the GPT Builder — showing its description, behavioral instructions, and conversation starters.

Final published version of the FocusMate AI assistant, fully configured and ready for user interaction.




Figure 21. Final FocusMate interface - the completed AI productivity coach designed to help users plan, focus, and stay motivated.

Link to chatbot: [FocusMate - ChatGPT Custom GPT](#)

Testing and Interaction Outputs

Output 1 - Motivation Scenario

When the user greeted the bot with “Hey!!!”, FocusMate responded enthusiastically, asking how the user was feeling and offering to plan the day together.

When the user replied, “*I’m tired and can’t focus, give me motivation,*” the chatbot responded empathetically:

“That sounds tough — when your energy’s low, it can feel impossible to get started. But remember: you don’t need a burst of motivation to begin — just one small action... You’re stronger than your tiredness — let’s build momentum, one gentle step at a time.”

The chatbot balanced empathy and practicality, showing how AI can motivate users through supportive, human-like language rather than commands.

Observation 1: The test confirmed that FocusMate reflects Design Thinking principles, leading with empathy and guiding users toward small, achievable actions.




Figure 22. FocusMate motivation response - demonstrating empathy-driven design where the AI offers gentle encouragement and practical actions to overcome fatigue and build momentum.

Output 2 - Planning Scenario

When prompted with “*I have 3 assignments due, help me plan today,*” FocusMate acknowledged the stress and asked for available time.

After learning that six hours were available, it created a balanced Pomodoro-based schedule:

- 10:00 – 12:00: Assignment 1 (deep focus)
- 12:00 – 12:30: Break/stretch / lunch
- 12:30 – 2:30: Assignment 2 (steady pace)
- 2:30 – 3:00: Snack/rest
- 3:00 – 5:00: Assignment 3 (final push)

It ended with a motivational check-in:

“If you use the Pomodoro method (25 min focus + 5 min break), you’ll stay refreshed throughout. How does that schedule feel — too tight, or just right for today’s energy level?”

Observation 2: The chatbot translated a vague problem into a structured, actionable plan while maintaining a calm and encouraging tone - demonstrating clarity, structure, and emotional awareness.




Figure 23. FocusMate planning response - showcasing the assistant’s structured Pomodoro-based schedule and supportive tone to help users manage multiple tasks with balance and motivation.

Reflection

Building FocusMate through the Design Thinking process deepened my understanding of how empathy enhances AI interactions. Testing and iteration revealed that small adjustments in tone or phrasing significantly affect user trust and engagement.

This project taught me that AI systems are most effective when they listen before advising, and when their design balances functionality with human warmth. FocusMate evolved from a task planner into a supportive digital companion that embodies ethical and empathetic AI principles.

AI Acknowledgment

ChatGPT-5 was used to assist in structuring, testing, and documenting the FocusMate chatbot. It supported the phrasing of dialogue examples, image captions, and the organization of layout in this report. All creative ideas, design decisions, testing observations, and final reflections were independently developed and verified by me.